

DATA ANALYSIS WITH PANDAS

This data was extracted from the Census Bureau Database found at

<http://www.census.gov/ftp/pub/DES/www/welcome.htm>

Donor: Ronny Kohavi and Barry Becker.

Extraction was done by Barry Becker from the 1994 Census database.

Exploratory analysis: Loading and exploring the dataset

```
In [1]: #Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data = pd.read_csv(r'C:\Users\LenovoX260\Desktop\Data Mining and Informatics Assignme
```

```
In [3]: data.head()
```

```
Out[3]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

`data.head()` - This code displays the first 5 rows of the dataset from 0-4

```
In [4]: data.head(2)
```

Out[4]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	ca
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	

Observation

data.head(2) - This shows the first 2 rows of the dataset from 0-1

In [5]: `data.head(10)`

Out[5]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female
5	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female
6	49	Private	160187	9th	5	Married-spouse-absent	Other-service	Not-in-family	Black	Female
7	52	Self-emp-not-inc	209642	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	Male
8	31	Private	45781	Masters	14	Never-married	Prof-specialty	Not-in-family	White	Female
9	42	Private	159449	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male

Observation

`data.head(10)` - This shows the first 10 rows of the data set from 0-9

```
In [6]: data.tail(2)
```

```
Out[6]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	M
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	F

Observation

`data.tail(2)` - This shows the last 2 rows of the data set

```
In [7]: data.shape
```

```
Out[7]: (32561, 15)
```

The above code shows the number of rows and columns in the dataset. In this data there are 32561 rows and 15 columns

Generating your unique dataset for this task

In this section we will be generating a unique dataset by replacing the last two digits in random state with 17

```
In [8]: data = data.sample(n=30000, random_state = 17)
```

```
In [9]: data.shape
```

```
Out[9]: (30000, 15)
```

Running the code `data.shape`. There are now 3000 rows and 15 columns in the dataset

```
In [10]: data.describe()
```

Out[10]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
count	30000.000000	3.000000e+04	30000.000000	30000.000000	30000.000000	30000.000000
mean	38.571033	1.896964e+05	10.084200	1076.818167	87.850200	40.471367
std	13.645176	1.055088e+05	2.572586	7412.566535	404.629371	12.388020
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.177670e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783410e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.372968e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

data.describe() - This code shows the descriptive statistics of all numerical attributes in the dataset

In [11]: `data['education-num'].value_counts()`

Out[11]:

9	9676
10	6706
13	4948
14	1589
11	1269
7	1069
12	986
6	852
4	602
15	532
5	477
8	410
16	382
3	303
2	153
1	46

Name: education-num, dtype: int64

The above code shows the occurrence of each values in the education_num attribute column sorted from the most to the least frequent in the dataset

In [12]: `data['education'].value_counts()`

```
Out[12]:
```

HS-grad	9676
Some-college	6706
Bachelors	4948
Masters	1589
Assoc-voc	1269
11th	1069
Assoc-acdm	986
10th	852
7th-8th	602
Prof-school	532
9th	477
12th	410
Doctorate	382
5th-6th	303
1st-4th	153
Preschool	46

```
Name: education, dtype: int64
```

The above code shows the occurrence of each values in the education attribute column sorted from the most to the least frequent in the dataset

```
In [13]: data = data.drop(['fnlwgt'], axis=1)
```

`data=data.drop` code is used to drop a column in the dataset. This was used to drop the 'fnlwgt' column which will result to the number of columns reducing from 15 columns from 14 columns

```
In [14]: data.shape
```

```
Out[14]: (30000, 14)
```

We run the `data.shape` code to confirm that the 'fnlwgt' column has been dropped off the dataset

```
In [15]: data.describe(include='all')
```

Out[15]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race
count	30000.000000	30000	30000	30000.000000	30000	30000	30000	30000
unique	NaN	9	16	NaN	7	15	6	5
top	NaN	Private	HS-grad	NaN	Married-civ-spouse	Prof-specialty	Husband	White
freq	NaN	20941	9676	NaN	13785	3812	12144	25659
mean	38.571033	NaN	NaN	10.084200	NaN	NaN	NaN	NaN
std	13.645176	NaN	NaN	2.572586	NaN	NaN	NaN	NaN
min	17.000000	NaN	NaN	1.000000	NaN	NaN	NaN	NaN
25%	28.000000	NaN	NaN	9.000000	NaN	NaN	NaN	NaN
50%	37.000000	NaN	NaN	10.000000	NaN	NaN	NaN	NaN
75%	48.000000	NaN	NaN	12.000000	NaN	NaN	NaN	NaN
max	90.000000	NaN	NaN	16.000000	NaN	NaN	NaN	NaN

The `data.describe(include='all')` code is used to show the descriptive statistics of all numerical and categorical data in the dataset.

```
In [16]: data['education'].value_counts()
```

```
Out[16]: HS-grad          9676
Some-college    6706
Bachelors       4948
Masters         1589
Assoc-voc       1269
11th            1069
Assoc-acdm       986
10th             852
7th-8th          602
Prof-school      532
9th              477
12th             410
Doctorate        382
5th-6th          303
1st-4th          153
Preschool         46
Name: education, dtype: int64
```

The above code shows the occurrence of each values in the education attribute column sorted from the most to the least frequent in the dataset

```
In [17]: data['education'].nunique()
```

```
Out[17]: 16
```

This code returns the count of unique values in the education column. There are 16 unique values in the education column.

```
In [18]: data['age'].value_counts()
```


```
Out[18]: 31    828
          34    815
          33    813
          23    812
          36    810
          ...
          83     5
          85     3
          88     2
          87     1
          86     1
          Name: age, Length: 73, dtype: int64
```

value_counts isn't an ideal way to generate the occurrence of each values in the age attribute column. To get a better view of age we will make use of a boxplot graph in visualizing age because age is a continuous value and is too lengthy to be analyzed using value_counts function

Lets visualize age using boxplot instead;


```
In [19]: data.boxplot(column='age')
```

```
Out[19]: <AxesSubplot:>
```




```
In [20]: data['age'].hist(bins=100)
```

```
Out[20]: <AxesSubplot:>
```



```
In [21]: data.age.hist(bins=100)
```

```
Out[21]: <AxesSubplot:>
```



```
In [22]: data['sex'].value_counts()
```



```
Out[22]: Male      20082
         Female    9918
         Name: sex, dtype: int64
```

`value_counts` by 'sex' is used to show the number of males and females that exist in the dataset. Here there are 20082 males and 9918 females in the dataset

```
In [23]: data.columns
```

```
Out[23]: Index(['age', 'workclass', 'education', 'education-num', 'marital-status',
               'occupation', 'relationship', 'race', 'sex', 'capital-gain',
               'capital-loss', 'hours-per-week', 'native-country', 'class-label'],
              dtype='object')
```

```
In [24]: data['workclass'].value_counts()
```

```
Out[24]: Private      20941
         Self-emp-not-inc  2353
         Local-gov      1920
         ?              1675
         State-gov      1178
         Self-emp-inc    1028
         Federal-gov     885
         Without-pay     14
         Never-worked     6
         Name: workclass, dtype: int64
```

`value_counts` by 'workclass' is used to show the number of the occurrence of each work sector workclasses attribute column sorted from the most to the least frequent in the dataset

```
In [25]: data['sex'].value_counts()
```

```
Out[25]: Male      20082
         Female    9918
         Name: sex, dtype: int64
```

The `value_count` by sex function above shows there are 20082 males and 9918 females in the dataset

Applying Groupby Functions in Order to Summarise the Data.

Average age of each gender in the given population

```
In [26]: data['age'].groupby([data['sex']]).mean()
```

```
Out[26]: sex
         Female    36.854104
         Male      39.418982
         Name: age, dtype: float64
```

Using `groupby` aggregate function we can derive the average age of each gender in the dataset.

In the above we grouped by 'sex' and calculated the average of 'age'. Hence the average age for Female is 36.854104 and Male is 39.418982

Average age of male and female across different education categories

```
In [27]: data['age'].groupby([data['sex'], data['education']]).mean()
```

```
Out[27]: sex      education
Female  10th      35.381818
        11th      30.687179
        12th      30.305970
        1st-4th   49.214286
        5th-6th   44.986486
        7th-8th   49.891156
        9th       42.044118
        Assoc-acdm 36.418848
        Assoc-voc  37.805252
        Bachelors  35.671096
        Doctorate  45.012500
        HS-grad    38.633834
        Masters    43.030488
        Preschool  42.933333
        Prof-school 40.011765
        Some-college 33.687621
Male    10th      38.150780
        11th      33.597938
        12th      33.246377
        1st-4th   45.090090
        5th-6th   41.851528
        7th-8th   48.292308
        9th       40.577713
        Assoc-acdm 37.865894
        Assoc-voc  38.948276
        Bachelors  40.284055
        Doctorate  48.311258
        HS-grad    39.101390
        Masters    44.515041
        Preschool  43.096774
        Prof-school 45.702461
        Some-college 36.995156
```

Name: age, dtype: float64

Using the groupby function we grouped by 'sex' and 'education' and computed the mean for 'age'

Average contribution to capital-gain of each sex and occupation category.

```
In [28]: data['capital-gain'].groupby([data['sex'], data['occupation']]).mean()
```

```
Out[28]: sex      occupation
        Female    ?              317.976864
          Adm-clerical      530.984985
          Craft-repair      815.635922
          Exec-managerial  1020.321731
          Farming-fishing   667.949153
          Handlers-cleaners 141.777070
          Machine-op-inspct 175.299413
          Other-service     165.004225
          Priv-house-serv   313.770992
          Prof-specialty    1212.759857
          Protective-serv   1525.614286
          Sales             228.426230
          Tech-support      662.320872
          Transport-moving  513.825000
        Male      ?              824.759690
          Adm-clerical      467.401918
          Armed-Forces       0.000000
          Craft-repair      627.810544
          Exec-managerial  2778.824294
          Farming-fishing   536.919240
          Handlers-cleaners 281.325088
          Machine-op-inspct 387.201800
          Other-service     232.912446
          Priv-house-serv   84.857143
          Prof-specialty    3548.676045
          Protective-serv   539.611006
          Sales             1911.073039
          Tech-support      691.472590
          Transport-moving  514.976412
```

Name: capital-gain, dtype: float64

Using the groupby function we grouped by 'sex' and 'occupation' and computed the mean for 'capital-gain'

Identifying the average capital-gain by males and females across different marital-status

```
In [29]: data['capital-gain'].groupby([data['sex'], data['marital-status']]).mean()
```

```
Out[29]: sex      marital-status
        Female    Divorced          452.100000
          Married-AF-spouse      204.076923
          Married-civ-spouse    1562.743119
          Married-spouse-absent  227.136842
          Never-married         342.128311
          Separated             194.290657
          Widowed              476.324573
        Male      Divorced          1093.146969
          Married-AF-spouse      912.250000
          Married-civ-spouse    1799.035076
          Married-spouse-absent  923.910448
          Never-married         422.352015
          Separated             789.347107
          Widowed              1005.943038
```

Name: capital-gain, dtype: float64

Maximum age across different races

```
In [30]: data['race'].value_counts()
```

```
Out[30]:
```

White	25659
Black	2842
Asian-Pac-Islander	966
Amer-Indian-Eskimo	290
Other	243

Name: race, dtype: int64

```
In [31]: data['age'].groupby([data['race']]).max()
```

```
Out[31]:
```

race	
Amer-Indian-Eskimo	80
Asian-Pac-Islander	90
Black	90
Other	77
White	90

Name: age, dtype: int64

90 is the maximum age amongst the Asian-Pac-Islander, Black and White races, using the groupby function above.

Are minimum and maximum age by sex same?

Minimum Age by Sex

```
In [32]: data['age'].groupby([data['sex']]).min()
```

```
Out[32]:
```

sex	
Female	17
Male	17

Name: age, dtype: int64

Minimum Age by Sex

```
In [33]: data['age'].groupby([data['sex']]).max()
```

```
Out[33]:
```

sex	
Female	90
Male	90

Name: age, dtype: int64

Using the groupby function the minimum and maximum age by sex are the same

The minimum age for Male and Female is 17 while the maximum age for Male and Female is 90 as well.

Data Visualization

```
In [34]: import matplotlib.pyplot as plt
         %matplotlib inline
```


```
In [35]: data.describe()
```

Out[35]:

	age	education-num	capital-gain	capital-loss	hours-per-week
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	38.571033	10.084200	1076.818167	87.850200	40.471367
std	13.645176	2.572586	7412.566535	404.629371	12.388020
min	17.000000	1.000000	0.000000	0.000000	1.000000
25%	28.000000	9.000000	0.000000	0.000000	40.000000
50%	37.000000	10.000000	0.000000	0.000000	40.000000
75%	48.000000	12.000000	0.000000	0.000000	45.000000
max	90.000000	16.000000	99999.000000	4356.000000	99.000000


In [36]: `data['age'].hist(bins=50)`

Out[36]: `<AxesSubplot:>`




In [37]: `data.boxplot(column='age')`

Out[37]: `<AxesSubplot:>`



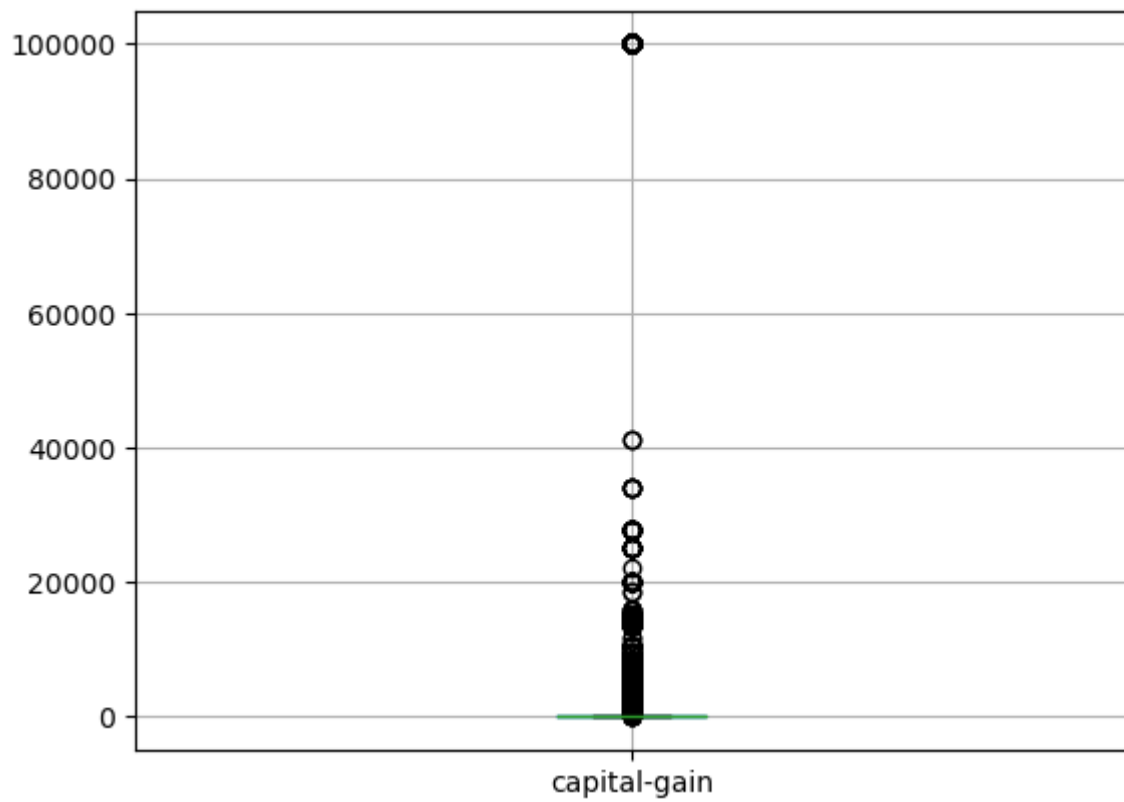
```
In [38]: data['capital-gain'].hist(bins=100)
```

```
Out[38]: <AxesSubplot:>
```




```
In [39]: data.boxplot(column='capital-gain')
```

```
Out[39]: <AxesSubplot:>
```



```
In [40]: data.boxplot(column='age', by = 'education', grid=False, rot = 45, fontsize = 10)
```

```
Out[40]: <AxesSubplot:title={'center':'age'}, xlabel='education'>
```




```
In [41]: data['education'].value_counts()
```

```
Out[41]: HS-grad          9676
Some-college       6706
Bachelors         4948
Masters           1589
Assoc-voc         1269
11th              1069
Assoc-acdm         986
10th              852
7th-8th           602
Prof-school       532
9th               477
12th              410
Doctorate         382
5th-6th           303
1st-4th           153
Preschool         46
Name: education, dtype: int64
```

```
In [42]: data.boxplot(column='capital-gain', by = 'education', grid=False, rot = 45, fontsize = 10)
```

```
Out[42]: <AxesSubplot:title={'center': 'capital-gain'}, xlabel='education'>
```

```
In [43]: data['marital-status'].value_counts()
```

```
Out[43]: Married-civ-spouse      13785
Never-married      9840
Divorced           4103
Separated           941
Widowed            919
Married-spouse-absent    391
Married-AF-spouse       21
Name: marital-status, dtype: int64
```

```
In [44]: data.apply(lambda x: sum(x.isnull()), axis = 0)
```

```
Out[44]: age                0
workclass                0
education                0
education-num            0
marital-status           0
occupation               0
relationship             0
race                     0
sex                      0
capital-gain             0
capital-loss             0
hours-per-week           0
native-country           0
class-label              0
dtype: int64
```

Data Transformation

```
In [45]: from sklearn.preprocessing import LabelEncoder
```

```
In [46]: data.head()
```

```
Out[46]:
```

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain
17121	54	Private	Some-college	10	Married-civ-spouse	Transport-moving	Husband	White	Male	0
363	43	Private	Bachelors	13	Divorced	Exec-managerial	Not-in-family	White	Male	0
17291	42	Self-emp-not-inc	HS-grad	9	Married-civ-spouse	Craft-repair	Husband	White	Male	0
16063	65	?	HS-grad	9	Married-civ-spouse	?	Husband	White	Male	0
25565	34	State-gov	Some-college	10	Married-spouse-absent	Adm-clerical	Unmarried	Asian-Pac-Islander	Female	0

```
In [47]: data.dtypes
```

```
Out[47]: age                int64
workclass            object
education            object
education-num        int64
marital-status       object
occupation           object
relationship         object
race                object
sex                 object
capital-gain         int64
capital-loss         int64
hours-per-week       int64
native-country       object
class-label          object
dtype: object
```

```
In [48]: columns = list(data.select_dtypes(exclude=['int64']))
```

```
In [49]: columns
```

```
Out[49]: ['workclass',
'education',
'marital-status',
'occupation',
'relationship',
'race',
'sex',
'native-country',
'class-label']
```

```
In [50]: data['class-label'].value_counts()
```

```
Out[50]: <=50K    22781
>50K        7219
Name: class-label, dtype: int64
```

```
In [51]: le = LabelEncoder()
for i in columns:
    #print(i)
    data[i] = le.fit_transform(data[i])
data.dtypes
```

```
Out[51]: age                int64
workclass            int32
education            int32
education-num        int64
marital-status       int32
occupation           int32
relationship         int32
race                int32
sex                 int32
capital-gain         int64
capital-loss         int64
hours-per-week       int64
native-country       int32
class-label          int32
dtype: object
```

```
In [52]: data.head()
```

Out[52]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain
17121	54	4	15	10	2	14	0	4	1	0
363	43	4	9	13	0	4	1	4	1	0
17291	42	6	11	9	2	3	0	4	1	0
16063	65	0	11	9	2	0	0	4	1	0
25565	34	7	15	10	3	1	4	1	0	0

In [53]: `data['workclass'].value_counts()`


Out[53]:

4	20941
6	2353
2	1920
0	1675
7	1178
5	1028
1	885
8	14
3	6

Name: workclass, dtype: int64

In [54]: `data['education'].hist(bins=100)`

Out[54]: <AxesSubplot:>



In [55]: `data.describe(include='all')`

Out[55]:

	age	workclass	education	education-num	marital-status	occupation	relations
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000
mean	38.571033	3.870767	10.299267	10.084200	2.611433	6.581300	1.448
std	13.645176	1.451644	3.864152	2.572586	1.507136	4.227661	1.607
min	17.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000
25%	28.000000	4.000000	9.000000	9.000000	2.000000	3.000000	0.000
50%	37.000000	4.000000	11.000000	10.000000	2.000000	7.000000	1.000
75%	48.000000	4.000000	12.000000	12.000000	4.000000	10.000000	3.000
max	90.000000	8.000000	15.000000	16.000000	6.000000	14.000000	5.000

Report

Summary of the outcome of data.describe()

In [56]: `data = pd.read_csv(r'C:\Users\LenovoX260\Desktop\Data Mining and Informatics Assignme`

In [57]: `data.describe()`

Out[57]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

The outcome of data.describe() function used here shows the descriptive statistics of all numerical attributes in the dataset

It shows the count, mean, standard deviation, minimum, first quartile, second quartile, third quartile and maximum value of age, education-num, capital-gain and hours-per-week attributes in the dataset.

The different data types (or attribute types) in data mining.

From the Adult dataset we have two types of attributes

1. Categorical or Qualitative Attribute.
2. Numerical or Quantitative Attribute.

Categorical or Qualitative Attributes -

These attribute takes qualitative values with qualitative characteristics. They are classified into two types;

- Nominal Attributes - This type of attribute provides enough information to differentiate between one object from another. Examples are; customer ID, student ID, zip codes, employee ID, gender, sex etc.
- Ordinal Attribute: The ordinal attribute value provides sufficient information to order the objects such as rankings, grades, street numbers, height etc

From the Adult dataset attributes like workclass, education, marital-status, occupation, relationship, race, sex and native-country are all classified as Categorical or Qualitative attributes.

Numerical or Quantitative Attributes -

This are data that contains whole numbers and decimals and have most properties of numbers. T Numerical attributes are further divided into;

- Binary Attribute: These are 0 and 1. Where 0 is the absence of any features and 1 is the inclusion of any characteristics. Numeric attribute: It is quantitative, such that quantity can be measured and represented in integer or real values, are of two types
- Interval Scaled attribute: It is measured on a scale of equal size units, these attributes allow us to compare values such as calendar dates temperature in Celsius or Fahrenheit.
- Ratio Scaled attribute: For ratio both differences and ratios are significant. For example, age, counts, temperature in Kelvin, length, and Weight.

From the Adult dataset we can see some Numerical or Quantitative attributes like age, education number, capital gain, capital loss and hours per week.

Country with the Highest Migrants.

Using the code below we can see that United-States have the largest number of migrants but in our answer we can't say the United-States have the largest migrant because the data was derived in the United States and the citizens of the United-States were included in the data which we obviously can't ignore in our analysis cause they are part of the dataset. Hence, in this case we would pick the country with the second largest number, people who are actually migrants in the United-States. Therefore, Mexico is the country with the largest number with 643 migrants after the United States.

```
In [58]: data = pd.read_csv(r'C:\Users\LenovoX260\Desktop\Data Mining and Informatics Assignme
```

```
In [61]: a=data["native-country"].value_counts()
a.head(2)
```

```
Out[61]: United-States    29170
Mexico                643
Name: native-country, dtype: int64
```

Occupation that represents more males than females

To get this we will apply the code below.

```
In [62]: data['sex'].groupby([data['occupation']]).value_counts()
```

```
Out[62]: occupation    sex
?                    Male    1002
                   Female    841
Adm-clerical         Female  2537
                   Male    1233
Armed-Forces         Male      9
Craft-repair         Male   3877
                   Female   222
Exec-managerial      Male   2907
                   Female  1159
Farming-fishing      Male    929
                   Female    65
Handlers-cleaners    Male   1206
                   Female   164
Machine-op-inspct    Male   1452
                   Female   550
Other-service        Female  1800
                   Male   1495
Priv-house-serv      Female   141
                   Male      8
Prof-specialty       Male   2625
                   Female  1515
Protective-serv      Male    573
                   Female    76
Sales                Male   2387
                   Female  1263
Tech-support         Male    580
                   Female   348
Transport-moving     Male   1507
                   Female    90
Name: sex, dtype: int64
```

From the code above, we can see there are 12 occupation that consist of more male than female. Occupation such as;

1. Craft repair - which has 3877 males than females.
2. Exec-managerial - consist of 2907 males and less females.
3. Farming-fishing - with 929 more males than females.
4. Handlers-cleaners - with 1206 males and females.
5. Machine-op-inspct - which has 1452 more males than females.
6. Prof-specialty - with 2625 more males than females.
7. Protective-serv - Has 573 more males than females.

8. Sales - Has 2387 more males than females.
9. Tech-support - Has 580 more males than females.
10. Transport-moving - Has more males of 1507 than females.

Finally, lets not forget the unidentified occupation and the armed-forces in the dataset contains more males and females. With the unidentified occupation having 1002 males and the armed forces having only just males with no females.

Difference between data.head() and data.tail()

In [63]: `data.head()`

Out[63]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

data.head() - This function is used to display the first 5 rows in the data set from 0-4

In [64]: `data.tail()`

Out[64]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	
32556	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Fer
32557	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	M
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Fer
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	M
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Fer



`data.tail()` - This code is used to display the last 5 rows in the data set.